

Exercise 1

We first load a dataset and examine its dimensions.

```
In [1]: # If you are running this on Google Colab, uncomment and run the following lines; o  
# from google.colab import drive  
# drive.mount('/content/drive')
```


```
In [1]: import math  
import numpy as np  
  
xy_data = np.load('Ex1_polyreg_data.npy')  
# If running on Google Colab change path to '/content/drive/MyDrive/IB-Data-Science  
  
np.shape(xy_data)
```

Out[1]: (70, 2)

The matrix `xy_data` contains 70 rows, each a data point of the form (x_i, y_i) for $i = 1, \dots, 70$.

1a) Plot the data in a scatterplot.

```
In [4]: import matplotlib.pyplot as plt  
  
# Split xy_data into x and y vectors  
x = xy_data[:, 0]  
y = xy_data[:, 1]  
  
plt.figure(figsize=(7, 4))  
plt.scatter(x, y, s=25) # s controls marker size  
plt.xlabel("x")  
plt.ylabel("y")  
plt.title("Scatter plot of (x, y) data")  
plt.grid(True, alpha=0.3)  
plt.tight_layout()  
plt.show()
```



1b) Write a function `polyreg` to fit a polynomial of a given order to a dataset.

The inputs to the function are a data matrix of dimension $N \times 2$, and $k \geq 0$, the order of the polynomial. The function should compute the coefficients of the polynomial $\beta_0 + \beta_1 x + \dots + \beta_k x^k$ via least squares regression, and should return the coefficient vector, the fit, and the vector of residuals.

If specified the degree k is greater than or equal to N , then the function must fit an order $(N - 1)$ polynomial and set the remaining coefficients to zero.

NOTE: You are *not* allowed to use the built-in function `np.polyfit`.

```
In [6]: import numpy as np

def polyreg(data_matrix, k):
    """
    Least-squares polynomial regression.

    Parameters
    -----
    data_matrix : (N, 2) ndarray
        First column is x_i, second column is y_i.
    k : int
        Requested polynomial degree (k >= 0).

    Returns
    -----
    beta : (k + 1,) ndarray
        Coefficient vector [ $\beta_0$ ,  $\beta_1$ , ...,  $\beta_k$ ]. If  $k \geq N$ , the last
        (k - N + 1) entries are zero-padded.
```

```

fit : (N,) ndarray
    The fitted y-values at the original x_i.
residuals : (N,) ndarray
    y_i - fit_i.
"""
x = data_matrix[:, 0]
y = data_matrix[:, 1]
N = x.size

# If k is too large, cap the effective degree at N-1
k_eff = min(k, N - 1)

# Design (Vandermonde) matrix with powers 0 ... k_eff
X = np.vander(x, N=k_eff + 1, increasing=True) # shape (N, k_eff+1)

# Solve the normal equations via Least-squares
beta_eff, *_ = np.linalg.lstsq(X, y, rcond=None) # shape (k_eff+1,)

# Pad with zeros if user asked for a higher degree than we could fit
if k_eff < k:
    beta = np.concatenate([beta_eff, np.zeros(k - k_eff)])
else:
    beta = beta_eff

# Compute fitted values and residuals
fit = X @ beta_eff # uses effective design matrix
residuals = y - fit

return beta, fit, residuals

```

Use the tests below to check the outputs of the function you have written:

```

In [7]: # Some tests to make sure your function is working correctly

xcol = np.arange(-1, 1.05, 0.1)
ycol = 2 - 7*xcol + 3*(xcol**2) # We are generating data according to y = 2 - 7x +
test_matrix = np.transpose(np.vstack((xcol,ycol)))
test_matrix.shape

beta_test = polyreg(test_matrix, k=2)[0]
assert((np.round(beta_test[0], 3) == 2) and (np.round(beta_test[1], 3) == -7) and (
# We want to check that using the function with k=2 recovers the coefficients exact

# Now check the zeroth order fit, i.e., the function gives the correct output with
beta_test = polyreg(test_matrix, k=0)[0]
res_test = polyreg(test_matrix, k=0)[2] #the last output of the function gives the

assert(np.round(beta_test, 3) == 3.1)
assert(np.round(np.linalg.norm(res_test), 3) == 19.937)

```

1c) Use `polyreg` to fit polynomial models for the data in `xy_data` for $k = 2, 3, 4$:

- Plot the fits for the three cases on the same plot together with the scatterplot of the data. The plots should be labelled and a legend included.
- Compute and print the SSE and R^2 coefficient for each of the three cases.
- Which of the three models you would choose? Briefly justify your choice.

```
In [ ]: import numpy as np
import matplotlib.pyplot as plt

# ----- helper: evaluate polynomial at arbitrary x -----
def eval_poly(beta, x_vals):
    """
    Evaluate the polynomial with coefficients beta ( $\beta_0, \beta_1, \dots$ ) at x_vals.
    """
    powers = np.arange(len(beta)) # exponents 0,1,2,...
    return np.sum(beta[np.newaxis, :] * x_vals[:, None]**powers, axis=1)

# ----- set-up -----
x = xy_data[:, 0]
y = xy_data[:, 1]
x_grid = np.linspace(x.min(), x.max(), 400) # smooth curve
degrees = [2, 3, 4]

plt.figure(figsize=(8, 5))
plt.scatter(x, y, s=25, label="data", zorder=3)

print("Degree |      SSE      |      R2")
print("-----")

best_k, best_r2 = None, -np.inf

for k in degrees:
    beta, y_fit, resid = polyreg(xy_data, k)

    # stats
    sse = np.sum(resid**2)
    ss_tot = np.sum((y - y.mean())**2)
    r2 = 1 - sse / ss_tot

    # track best
    if r2 > best_r2:
        best_k, best_r2 = k, r2

    # plot
    plt.plot(x_grid,
             eval_poly(beta, x_grid),
             label=f"degree {k}",
             linewidth=2)


    print(f"{k:>6} | {sse:11.4f} | {r2:6.4f}")

plt.xlabel("x")
plt.ylabel("y")
plt.title("Polynomial fits to data")
plt.legend()
```

```
plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()

print(f"\nChosen model: degree {best_k}")
```

Degree	SSE	R ²
2	172.1810	0.8876
3	152.4058	0.9005
4	151.2278	0.9013



Chosen model: degree 4

State which model you choose and briefly justify your choice.

-----#

Model-selection rationale:

SSE drops and R^2 rises noticeably when the degree increases from $k = 2$ to 3.

Going from $k = 3$ to $k = 4$ yields only a tiny gain in R^2 (diminishing returns

for an extra parameter) and risks fitting noise. Therefore, degree-3 is the

sweet-spot: it captures the systematic variation without unnecessary

complexity.

-----#

1d) For the model you have chosen in the previous part (either $k = 2/3/4$):

- Plot the residuals in a scatter plot.
- Plot a histogram of the residuals along with a Gaussian pdf with zero mean and the same standard deviation as the residuals.

```
In [9]: import numpy as np
import matplotlib.pyplot as plt

# --- choose the degree -----
k = best_k          # or set k = 3 (for example) if you know your choice

# --- refit and obtain residuals -----
beta, y_fit, resid = polyreg(xy_data, k)
x = xy_data[:, 0]

# ----- 1) residuals vs. x scatter plot -----
plt.figure(figsize=(7, 4))
plt.scatter(x, resid, s=25)
plt.axhline(0, color='k', linewidth=1)
plt.xlabel("x")
plt.ylabel("Residual")
plt.title(f"Residuals vs. x (degree {k})")
plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()

# ----- 2) histogram + Gaussian pdf -----
sigma = np.std(resid, ddof=1)
z = np.linspace(resid.min(), resid.max(), 400)
gauss_pdf = (1 / (np.sqrt(2 * np.pi) * sigma)) * np.exp(-0.5 * (z / sigma) ** 2)
```

```

plt.figure(figsize=(7, 4))
plt.hist(resid, bins="auto", density=True, alpha=0.6, edgecolor='k', label="Residuals")
plt.plot(z, gauss_pdf, linewidth=2, label=f"N(0, {sigma:.2f}^2) PDF")
plt.xlabel("Residual value")
plt.ylabel("Density")
plt.title(f"Histogram of residuals (degree {k}) with Gaussian overlay")
plt.legend()
plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()

```

